

Forecasting ARIMA(1,1,1) Series

ARIMA(1,1,1)

1. We generate the data assuming the true process is known. Then we can compare the estimation result to the truth to ensure the coding is right.
2. In general, an ARIMA(1,1,1) process is

$$\Delta y_t = d + \eta_t \quad (1)$$

$$\eta_t = \phi_1 \eta_{t-1} + e_t + \theta_1 e_{t-1} \quad (2)$$

In words, the first difference Δy_t is a zero-mean ARMA(1,1) process η_t plus the drift term d .

3. By substituting $\eta_t = y_t - y_{t-1} - d$, the same ARIMA(1,1,1) process can be written as

$$(y_t - y_{t-1} - d) = \phi_1 (y_{t-1} - y_{t-2} - d) + e_t + \theta_1 e_{t-1} \quad (3)$$

where d is the drift term; ϕ_1 is the AR coefficient; θ_1 is the MA coefficient.

4. Here we let $d = 0.2$, $\phi_1 = 0.7$, $\theta_1 = -0.5$. Notice that the nonzero drift term causes the series to be trending.

Generating ARIMA(1,1,1)

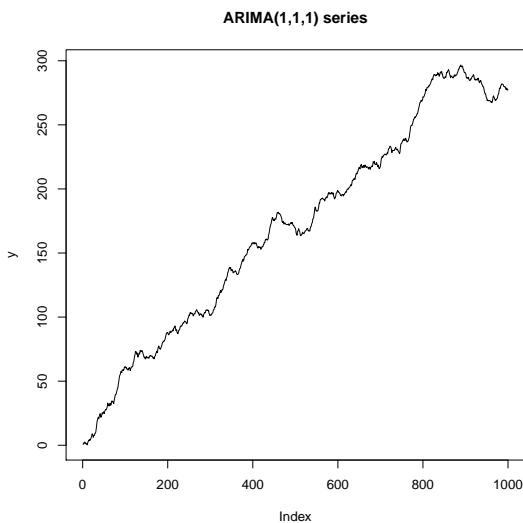
We use R loop to generate η_t and then y_t

```
set.seed(12345)
T = 1000
tr = 1:T
e = rnorm(T)
dy = rep(0, T)
y = rep(0, T)
dy[1] = e[1]
y[1] = e[1]
for (t in 2:T) {
  dy[t] = 0.7*dy[t-1]+e[t]-0.5*e[t-1]
  y[t] = 0.2 + y[t-1]+dy[t]
}
```

In the codes, the ARMA(1,1) η_t is denoted as $dy[t]$

Plotting ARIMA(1,1,1)

The series has an upward trend due to the positive drift term $d = 0.2$. The trend, along with the smoothness, signifies nonstationarity.



Deterministic and Stochastic Trends

We can show that the ARIMA(1,1,1) process is trending

$$y_t = y_0 + dt + (\eta_1 + \eta_2 + \dots + \eta_t)$$

There is a global deterministic trend dt if $d \neq 0$. Even if $d = 0$, there can be local stochastic trend $(\eta_1 + \eta_2 + \dots + \eta_t)$. In graph, the deterministic trend can be dominating. We can estimate the drift term, which is the mean value of Δy_t , without running regression

```
mean(d.y, na.rm=T)  
[1] 0.2769492
```

Estimating ARIMA(1,1,1)

1. Estimating ARIMA(1,1,1) for y_t is the same as estimating ARIMA(1,0,1) for Δy_t
2. However, DO NOT use `arima(y, order = c(1,1,1))` because this assumes zero drift term!!
3. Instead, use `arima(d.y, order = c(1,0,1))`. But be careful, the reported intercept actually is the drift term

```
arima(x = d.y, order = c(1, 0, 1))
```

Coefficients:

	ar1	ma1	intercept
	0.7246	-0.5197	0.2767
s.e.	0.0698	0.0871	0.0550

log likelihood = -1416.09, aic = 2840.17

The estimated $\hat{\phi}_1 = 0.7246$, $\hat{\theta}_1 = -0.5197$, $\hat{d} = 0.2767$ are all close to the true values, and are significant.

Estimating Error Terms

The error term e_t in (2) is unobservable. According to (3), we can show

$$\begin{aligned}e_t &= (y_t - y_{t-1} - d) - \phi_1(y_{t-1} - y_{t-2} - d) - \theta_1 e_{t-1} \\&= y_t - (1 - \phi_1)d - (1 + \phi_1)y_{t-1} + \phi_1 y_{t-2} - \theta_1 e_{t-1}\end{aligned}$$

So we use the codes below to estimate e_1, e_2, \dots, e_t in a recursive way

```
ehat = rep(0, T)
ehat[1] = y[1]
ehat[2] = y[2] - (1-phi1)*dhat - (1+phi1)*y[1] - theta1*ehat[1]
for (t in 3:T) ehat[t] = y[t] - (1-phi1)*dhat - (1+phi1)*y[t-1] + phi1*
```

Built-in Function for Estimating Error Terms

Except the early observations, our estimated error terms are very close to the ones reported by R built-in function `resid`

```
cbind(resid(arima(d.y, order = c(1,0,1))) [1:10], ehat[1:10])
```

	[,1]	[,2]
[1,]	NA	0.5855288
[2,]	0.7187304	0.8303763
[3,]	-0.1458968	-0.0739732
[4,]	-0.4963657	-0.4612424
[5,]	0.5616954	0.5796414
[6,]	-1.8649630	-1.8558010
[7,]	0.5906055	0.5951970
[8,]	-0.3141356	-0.3117367
[9,]	-0.3219821	-0.3207393
[10,]	-0.9561814	-0.9555372

Forecasting ARIMA(1,1,1)

1. Assuming our sample ends at T .
2. Rewrite (3) as

$$y_t = (1 - \phi_1)d + (1 + \phi_1)y_{t-1} - \phi_1 y_{t-2} + e_t + \theta_1 e_{t-1} \quad (4)$$

3. The one-step and two-step (out-of-sample) forecasts are

$$E(y_{T+1}|\Omega_T) = (1 - \phi_1)d + (1 + \phi_1)y_T - \phi_1 y_{T-1} + \theta_1 e_T \quad (5)$$

$$E(y_{T+2}|\Omega_T) = (1 - \phi_1)d + (1 + \phi_1)E(y_{T+1}|\Omega_T) - \phi_1 y_T \quad (6)$$

For k -th horizon ($k > 2$) we have

$$E(y_{T+k}|\Omega_T) = (1 - \phi_1)d + (1 + \phi_1)E(y_{T+k-1}|\Omega_T) - \phi_1 E(y_{T+k-2}|\Omega_T) \quad (7)$$

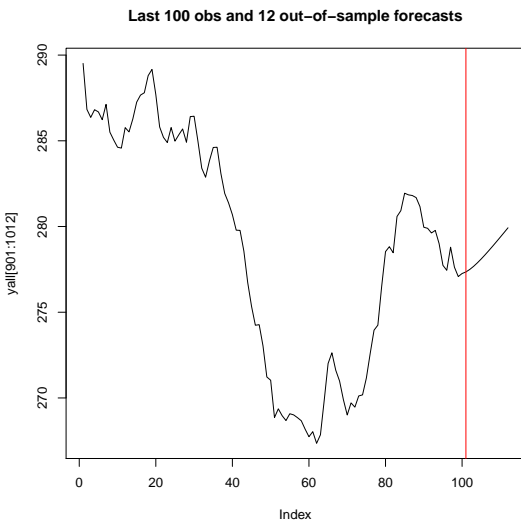
R Codes

Again, we use R loop to generate 12 out-of-sample forecasts, and display the first 5 forecasts

```
f = rep(0, 12)
f[1] = (1-phi1)*dhat+(1+phi1)*y[T]-phi1*y[T-1]+theta1*ehat[T]
f[2] = (1-phi1)*dhat+(1+phi1)*f[1]-phi1*y[T]
for (t in 3:12) f[t] = (1-phi1)*dhat+(1+phi1)*f[t-1]-phi1*f[t-2]
f[1:5]
[1] 277.3508 277.4945 277.6748 277.8817 278.1078
```


Plotting Forecasting Values I

```
yall = c(y, f)
plot(yall[901:1012], type="l", main="Last 100 obs and 12 out-of-samp
abline(v = c(101), col = "red", lty=1)
```



Plotting Forecasting Values II

```
plot(c(yall[901:1000], rep(NA, 12)), type="l", lwd=2, lty=1, main="L  
lines(c(rep(NA, 100), f), lty=2, col = "red", lwd=2)
```



R Forecast Package

There is a Forecast package that can provide the out-of-sample forecasts

```
library(forecast)
```

```
forecast(Arima(y,order=c(1,1,1),include.drift=T),h=12)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
1001	277.3508	276.0693	278.6324	275.3908	279.3108
1002	277.4945	275.4878	279.5012	274.4255	280.5635
1003	277.6748	275.0224	280.3272	273.6182	281.7313
1004	277.8816	274.6349	281.1283	272.9162	282.8471
1005	278.1077	274.3089	281.9066	272.2979	283.9175
1006	278.3478	274.0336	282.6619	271.7498	284.9457
1007	278.5979	273.8010	283.3948	271.2617	285.9341
1008	278.8554	273.6047	284.1060	270.8252	286.8855
1009	279.1181	273.4395	284.7967	270.4335	287.8028
1010	279.3847	273.3011	285.4684	270.0806	288.6889

Remarks

1. We obtain the same forecasts!
2. Arima is the function in the forecast package, which is different from arima in the stats package.
3. Notice that `include.drift=T` allows for a nonzero drift term.