

# Forecasting time series using R

**Professor Rob J Hyndman**

27 October 2011



**MONASH** University

# Outline

- 1 **Time series in R**
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# Australian GDP

```
ausgdp <- ts(scan("gdp.dat"), frequency=4,  
              start=1971+2/4)
```

# Australian GDP

```
ausgdp <- ts(scan("gdp.dat"), frequency=4,  
             start=1971+2/4)
```

- Class: `ts`
- Print and plotting methods available.

# Australian GDP

```
ausgdp <- ts(scan("gdp.dat"), frequency=4,  
              start=1971+2/4)
```


- Class: `ts`
- Print and plotting methods available.

```
> ausgdp
```

	Qtr1	Qtr2	Qtr3	Qtr4
1971			4612	4651
1972	4645	4615	4645	4722
1973	4780	4830	4887	4933
1974	4921	4875	4867	4905
1975	4938	4934	4942	4979
1976	5028	5079	5112	5127
1977	5130	5101	5072	5069

# Australian GDP

```
> plot(ausgdp)
```



# Residential electricity sales

```
> elecsales
```

```
Time Series:
```

```
Start = 1989
```

```
End = 2008
```

```
Frequency = 1
```

```
[1] 2354.34 2379.71 2318.52 2468.99 2386.09 2569.47  
[7] 2575.72 2762.72 2844.50 3000.70 3108.10 3357.50  
[13] 3075.70 3180.60 3221.60 3176.20 3430.60 3527.48  
[19] 3637.89 3655.00
```

# Useful packages

**Time series task view:** <http://cran.r-project.org/web/views/TimeSeries.html>



# Useful packages

**Time series task view:** <http://cran.r-project.org/web/views/TimeSeries.html>

---

<b>forecast</b>	for forecasting functions
<b>tseries</b>	for unit root tests and GARCH models
<b>Mcomp</b>	for the M-competition and M3-competition data
<b>fma</b>	for data from Makridakis, Wheelwright & Hyndman (1998)
<b>expsmooth</b>	for data from Hyndman et al. (2008)
<b>fpp</b>	for data from Hyndman & Athanasopoulos (forthcoming).

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods**
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# Some simple forecasting methods

## Mean method

- Forecast of all future values is equal to mean of historical data  $\{y_1, \dots, y_n\}$ .

# Some simple forecasting methods

## Mean method

- Forecast of all future values is equal to mean of historical data  $\{y_1, \dots, y_n\}$ .

- Forecasts:

$$\hat{y}_{n+h|n} = \bar{y} = (y_1 + \dots + y_n)/n$$

# Some simple forecasting methods

## Naïve method

- Forecasts equal to last observed value.

# Some simple forecasting methods

## Naïve method

- Forecasts equal to last observed value.
- Forecasts:  $\hat{y}_{n+h|n} = y_n$ .

# Some simple forecasting methods

## Naïve method

- Forecasts equal to last observed value.
- Forecasts:  $\hat{y}_{n+h|n} = y_n$ .
- Optimal for efficient stock markets.

# Some simple forecasting methods

## Seasonal naïve method

- Forecasts equal to last value from same season.



# Some simple forecasting methods

## Seasonal naïve method

- Forecasts equal to last value from same season.
- Forecasts:  $\hat{y}_{n+h|n} = y_{n-m}$  where  $m = \text{seasonal period}$  and  $k = \lfloor (h - 1)/m \rfloor + 1$ .

# Some simple forecasting methods

## Drift method

- Forecasts equal to last value plus average change.

# Some simple forecasting methods

## Drift method

- Forecasts equal to last value plus average change.
- Forecasts:

$$\begin{aligned}\hat{y}_{n+h|n} &= y_n + \frac{h}{n-1} \sum_{t=2}^n (y_t - y_{t-1}) \\ &= y_n + \frac{h}{n-1} (y_n - y_1).\end{aligned}$$

# Some simple forecasting methods

## Drift method

- Forecasts equal to last value plus average change.
- Forecasts:

$$\begin{aligned}\hat{y}_{n+h|n} &= y_n + \frac{h}{n-1} \sum_{t=2}^n (y_t - y_{t-1}) \\ &= y_n + \frac{h}{n-1} (y_n - y_1).\end{aligned}$$

- Equivalent to extrapolating a line drawn between first and last observations.


# Some simple forecasting methods

Australian annual beer production



# Some simple forecasting methods

**Australian annual beer production**



# Some simple forecasting methods

- Mean: `meanf(x, h=20)`

# Some simple forecasting methods

- Mean: `meanf(x, h=20)`
- Naive: `naive(x, h=20)` or `rwf(x, h=20)`



# Some simple forecasting methods

- Mean: `meanf(x, h=20)`
- Naive: `naive(x, h=20)` or `rwf(x, h=20)`
- Seasonal naive: `snaive(x, h=20)`

# Some simple forecasting methods

- Mean: `meanf(x,h=20)`
- Naive: `naive(x,h=20)` or `rwf(x,h=20)`
- Seasonal naive: `snaive(x,h=20)`
- Drift: `rwf(x,drift=TRUE,h=20)`

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

## forecast class contains

- Original series

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

## forecast class contains

- Original series
- **Point forecasts**

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

## forecast class contains

- Original series
- Point forecasts
- Prediction interval

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

## forecast class contains

- Original series
- Point forecasts
- Prediction interval
- **Forecasting method used**

# forecast objects in R

Functions that output a **forecast** object:

- `meanf()`
- `naive()`, `snaive()`
- `rwf()`
- `croston()`
- `stlf()`
- `ses()`
- `holt()`, `hw()`
- `splinef`
- `thetaf`
- `forecast()`

## forecast class contains

- Original series
- Point forecasts
- Prediction interval
- Forecasting method used
- **Residuals and in-sample one-step forecasts**



# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy**
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# Measures of forecast accuracy

Let  $y_t$  denote the  $t$ th observation and  $f_t$  denote its forecast, where  $t = 1, \dots, n$ . Then the following measures are useful.

$$\text{MAE} = n^{-1} \sum_{t=1}^n |y_t - f_t|$$

$$\text{MSE} = n^{-1} \sum_{t=1}^n (y_t - f_t)^2 \quad \text{RMSE} = \sqrt{n^{-1} \sum_{t=1}^n (y_t - f_t)^2}$$

$$\text{MAPE} = 100n^{-1} \sum_{t=1}^n |y_t - f_t| / |y_t|$$

# Measures of forecast accuracy

Let  $y_t$  denote the  $t$ th observation and  $f_t$  denote its forecast, where  $t = 1, \dots, n$ . Then the following measures are useful.

$$\text{MAE} = n^{-1} \sum_{t=1}^n |y_t - f_t|$$

$$\text{MSE} = n^{-1} \sum_{t=1}^n (y_t - f_t)^2 \quad \text{RMSE} = \sqrt{n^{-1} \sum_{t=1}^n (y_t - f_t)^2}$$

$$\text{MAPE} = 100n^{-1} \sum_{t=1}^n |y_t - f_t| / |y_t|$$

- MAE, MSE, RMSE are all scale dependent.

# Measures of forecast accuracy

Let  $y_t$  denote the  $t$ th observation and  $f_t$  denote its forecast, where  $t = 1, \dots, n$ . Then the following measures are useful.

$$\text{MAE} = n^{-1} \sum_{t=1}^n |y_t - f_t|$$

$$\text{MSE} = n^{-1} \sum_{t=1}^n (y_t - f_t)^2 \quad \text{RMSE} = \sqrt{n^{-1} \sum_{t=1}^n (y_t - f_t)^2}$$

$$\text{MAPE} = 100n^{-1} \sum_{t=1}^n |y_t - f_t| / |y_t|$$

- MAE, MSE, RMSE are all scale dependent.
- MAPE is scale independent but is only sensible if  $y_t \gg 0$  for all  $i$ , and  $y$  has a natural zero.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = n^{-1} \sum_{t=1}^n |y_t - f_t|/q$$

where  $q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = n^{-1} \sum_{t=1}^n |y_t - f_t|/q$$

where  $q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

Proposed by Hyndman and Koehler (IJF, 2006)

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = n^{-1} \sum_{t=1}^n |y_t - f_t| / q$$

where  $q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

For non-seasonal time series,

$$q = (n - 1)^{-1} \sum_{t=2}^n |y_t - y_{t-1}|$$

works well. Then MASE is equivalent to MAE relative to a naive method.

# Measures of forecast accuracy

## Mean Absolute Scaled Error

$$\text{MASE} = n^{-1} \sum_{t=1}^n |y_t - f_t|/q$$

where  $q$  is a stable measure of the scale of the time series  $\{y_t\}$ .

For seasonal time series,

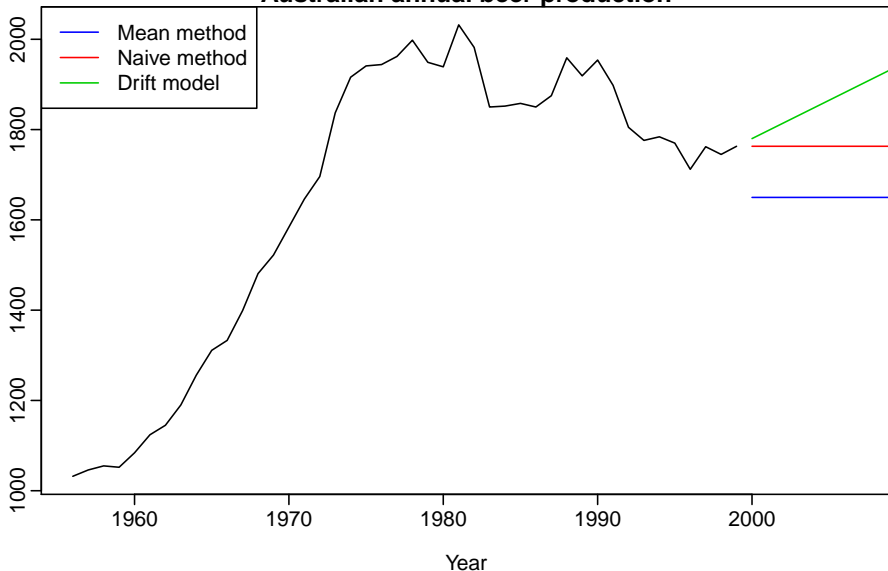
$$q = (n - m)^{-1} \sum_{t=m+1}^n |y_t - y_{t-m}|$$

works well. Then MASE is equivalent to MAE relative to a seasonal naive method.




# Measures of forecast accuracy

Australian annual beer production



# Measures of forecast accuracy

Australian annual beer production



# Measures of forecast accuracy

## Mean method

RMSE	MAE	MAPE	MASE
72.4223	68.6477	3.9775	1.5965

## Naïve method

RMSE	MAE	MAPE	MASE
50.2382	44.6250	2.6156	1.0378

## Drift method

RMSE	MAE	MAPE	MASE
134.6788	121.1250	7.0924	2.8169

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing**
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# Exponential smoothing


## Classic Reference



Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, 3rd ed., Wiley: NY.

# Exponential smoothing

## Classic Reference



Makridakis, Wheelwright and Hyndman (1998) *Forecasting: methods and applications*, 3rd ed., Wiley: NY.

## Current Reference



Hyndman, Koehler, Ord and Snyder (2008) *Forecasting with exponential smoothing: the state space approach*, Springer-Verlag: Berlin.

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	<b>N,N</b>	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

(N,N): Simple exponential smoothing



# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	<b>A,N</b>	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**(N,N):** Simple exponential smoothing

**(A,N):** Holt's linear method

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	<b>A,A</b>	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**(N,N):** Simple exponential smoothing

**(A,N):** Holt's linear method

**(A,A):** Additive Holt-Winters' method

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	<b>A,M</b>
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**(N,N):** Simple exponential smoothing

**(A,N):** Holt's linear method

**(A,A):** Additive Holt-Winters' method

**(A,M):** Multiplicative Holt-Winters' method

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	<b>A<sub>d</sub>,M</b>
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**(N,N):** Simple exponential smoothing

**(A,N):** Holt's linear method

**(A,A):** Additive Holt-Winters' method

**(A,M):** Multiplicative Holt-Winters' method

**(A<sub>d</sub>,M):** Damped multiplicative Holt-Winters' method

# Exponential smoothing methods

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

(N,N): Simple exponential smoothing

(A,N): Holt's linear method

(A,A): Additive Holt-Winters' method

(A,M): Multiplicative Holt-Winters' method

(A<sub>d</sub>,M): Damped multiplicative Holt-Winters' method

**There are 15 separate exponential smoothing methods.**

# R functions


- `HoltWinters()` implements methods (N,N), (A,N), (A,A) and (A,M). Initial states are selected heuristically. Smoothing parameters are estimated by minimizing MSE.

# R functions

- `HoltWinters()` implements methods (N,N), (A,N), (A,A) and (A,M). Initial states are selected heuristically. Smoothing parameters are estimated by minimizing MSE.
- Use `ets(x,model="ZMA",damped=FALSE)` to estimate parameters for the (M,A) method. All other methods similarly. Initial states and smoothing parameters are estimated using maximum likelihood estimation (see later).

# Exponential smoothing


Forecasts from ETS models






# Exponential smoothing

Forecasts from ETS models




# Exponential smoothing

Forecasts from ETS models




# Exponential smoothing

Forecasts from ETS models



# Exponential smoothing

Forecasts from ETS models



# Measures of forecast accuracy

## Mean method

RMSE	MAE	MAPE	MASE
72.4223	68.6477	3.9775	1.5965

## Naïve method


RMSE	MAE	MAPE	MASE
50.2382	44.6250	2.6156	1.0378

## Auto ETS model

RMSE	MAE	MAPE	MASE
27.3974	22.4014	1.3150	0.5210

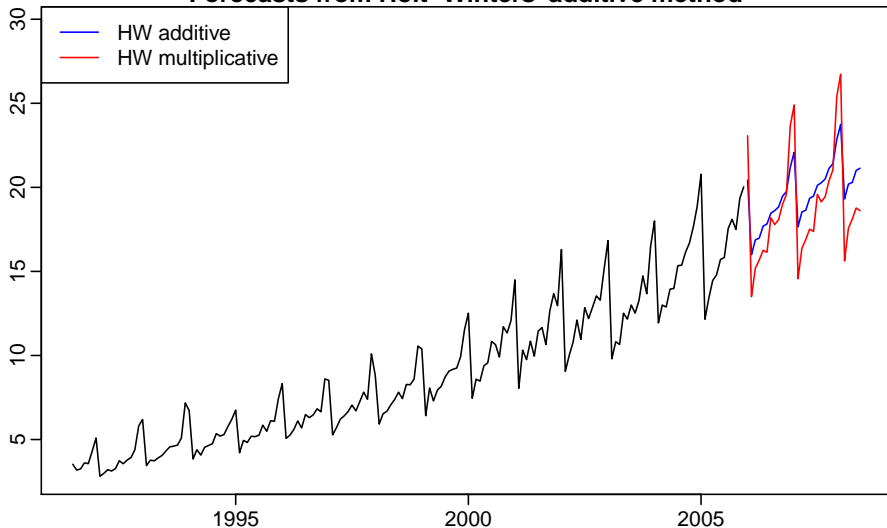
# Exponential smoothing

Forecasts from Holt-Winters' additive method




# Exponential smoothing

Forecasts from Holt-Winters' additive method



# Exponential smoothing

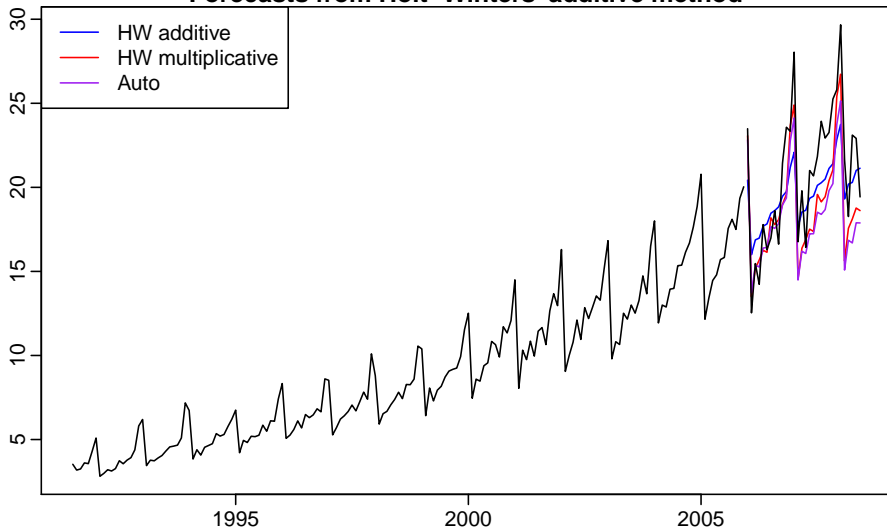
Forecasts from Holt-Winters' additive method





# Exponential smoothing

Forecasts from Holt-Winters' additive method



# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error,Trend,Seasonal*)

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error,Trend,Seasonal*)  
**Exponential Smoothing**

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	<b>N,N</b>	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error,Trend,Seasonal*)  
**Exponential Smoothing**

**ETS(A,N,N):** Simple exponential smoothing with additive errors

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	<b>A,N</b>	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error, Trend, Seasonal*)  
**Exponential Smoothing**

**ETS(A,A,N):** Holt's linear method with additive errors

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	<b>A,A</b>	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error, Trend, Seasonal*)  
**Exponential Smoothing**

**ETS(A,A,A):** Additive Holt-Winters' method with additive errors

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	<b>A,M</b>
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error, Trend, Seasonal*)  
**Exponential Smoothing**

**ETS(M,A,M):** Multiplicative Holt-Winters' method with multiplicative errors



# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	<b>A<sub>d</sub>,N</b>	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error,Trend,Seasonal*)  
**Exponential Smoothing**

**ETS(A,A<sub>d</sub>,N)**: Damped trend method with additive errors

# Exponential smoothing

Trend Component		Seasonal Component		
		N (None)	A (Additive)	M (Multiplicative)
N	(None)	N,N	N,A	N,M
A	(Additive)	A,N	A,A	A,M
A <sub>d</sub>	(Additive damped)	A <sub>d</sub> ,N	A <sub>d</sub> ,A	A <sub>d</sub> ,M
M	(Multiplicative)	M,N	M,A	M,M
M <sub>d</sub>	(Multiplicative damped)	M <sub>d</sub> ,N	M <sub>d</sub> ,A	M <sub>d</sub> ,M

**General notation** **ETS**(*Error,Trend,Seasonal*)  
**Exponential Smoothing**

**There are 30 separate models in the ETS framework**

# Exponential smoothing

```
fit <- ets(a10train)
fit2 <- ets(a10train,model="MMM",damped=FALSE)
fcast1 <- forecast(fit, h=30)
fcast2 <- forecast(fit2, h=30)
```

# Exponential smoothing

```
fit <- ets(a10train)
fit2 <- ets(a10train,model="MMM",damped=FALSE)
fcast1 <- forecast(fit, h=30)
fcast2 <- forecast(fit2, h=30)
```

```
ets(y, model="ZZZ", damped=NULL, alpha=NULL,
    beta=NULL, gamma=NULL, phi=NULL,
    additive.only=FALSE,
    lower=c(rep(0.0001,3),0.80),
    upper=c(rep(0.9999,3),0.98),
    opt.crit=c("lik","amse","mse","sigma"), nmse=3,
    bounds=c("both","usual","admissible"),
    ic=c("aic","aicc","bic"), restrict=TRUE)
```

# Exponential smoothing

```
> fit  
ETS(M,Ad,M)
```

```
Smoothing parameters:
```

```
alpha = 0.303  
beta  = 0.0205  
gamma = 1e-04  
phi   = 0.9797
```

```
Initial states:
```

```
l = 3.2875  
b = 0.0572  
s=0.9224 0.9252 0.8657 0.8763 0.7866 1.3149  
      1.2457 1.0641 1.0456 0.989 0.9772 0.9874
```

```
sigma: 0.0541
```

```
          AIC      AICc      BIC  
636.3248 640.2479 690.0287
```

# Exponential smoothing

```
> fit2  
ETS(M,M,M)
```

Smoothing parameters:

alpha = 0.3899

beta = 0.0766

gamma = 1e-04

Initial states:

l = 3.3312

b = 1.0057

s=0.923 0.9302 0.8569 0.8681 0.7796 1.334

1.2548 1.0505 1.0392 0.9867 0.9783 0.9988

sigma: 0.0565

AIC	AICc	BIC
652.3502	655.8151	702.8951

# Forecast accuracy

```
> accuracy(fcast1,a10test)
      RMSE      MAE      MAPE      MASE
3.3031418  2.7067125  12.3566798  2.6505174

> accuracy(fcast2,a10test)
      RMSE      MAE      MAPE      MASE
3.0001797  2.4143619  11.1033867  2.3642363
```

# Exponential smoothing

## ets() function

- Automatically chooses a model by default using the AIC, AICc or BIC.



# Exponential smoothing

## ets() function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping

# Exponential smoothing

## `ets()` function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model

# Exponential smoothing

## ets() function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)

# Exponential smoothing

## `ets()` function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)
- Produces an object of class `ets`.

# Exponential smoothing

## `ets()` function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)
- Produces an object of class `ets`.

# Exponential smoothing

## `ets()` function

- Automatically chooses a model by default using the AIC, AICc or BIC.
- Can handle any combination of trend, seasonality and damping
- Produces prediction intervals for every model
- Ensures the parameters are admissible (equivalent to invertible)
- Produces an object of class `ets`.

Automatic ETS algorithm due to Hyndman et al (IJF, 2002), updated in Hyndman et al (2008).

# Exponential smoothing

## ets objects

- **Methods:** `coef()`, `plot()`,  
`summary()`, `residuals()`,  
`fitted()`, `simulate()` and  
`forecast()`

# Exponential smoothing

## ets objects

- **Methods:** `coef()`, `plot()`, `summary()`, `residuals()`, `fitted()`, `simulate()` and `forecast()`
- `plot()` function shows time plots of the original time series along with the extracted components (level, growth and seasonal).



# Exponential smoothing

`ets()` function also allows refitting model to new data set.

```
> fit <- ets(a10train)

> test <- ets(a10test, model = fit)

> accuracy(test)
      RMSE      MAE      MAPE      MASE
1.51968  1.28391  6.50663  0.41450

> accuracy(forecast(fit,30), a10test)
      RMSE      MAE      MAPE      MASE
3.30314  2.70672 12.35668  2.65052
```

# Automatic forecasting

## Why use an automatic procedure?

- 1 Most users are not very expert at fitting time series models.

# Automatic forecasting

## Why use an automatic procedure?

- 1 Most users are not very expert at fitting time series models.
- 2 Most experts cannot beat the best automatic algorithms.

# Automatic forecasting

## Why use an automatic procedure?

- 1 Most users are not very expert at fitting time series models.
- 2 Most experts cannot beat the best automatic algorithms.
- 3 Many businesses and industries need thousands of forecasts every week/month.

# Automatic forecasting

## Why use an automatic procedure?

- 1 Most users are not very expert at fitting time series models.
- 2 Most experts cannot beat the best automatic algorithms.
- 3 Many businesses and industries need thousands of forecasts every week/month.
- 4 Some multivariate forecasting methods depend on many univariate forecasts.

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations**
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .



# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- $\lambda = 1$ : (No substantive transformation)

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- $\lambda = 1$ : (No substantive transformation)
- $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- $\lambda = 1$ : (No substantive transformation)
- $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)
- $\lambda = 0$ : (Natural logarithm)

# Transformations to stabilize the variance

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (y_t^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- $\lambda = 1$ : (No substantive transformation)
- $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)
- $\lambda = 0$ : (Natural logarithm)
- $\lambda = -1$ : (Inverse plus 1)

# Box-Cox transformations

# Back-transformation

We must reverse the transformation (or *back-transform*) to obtain forecasts on the original scale. The reverse Box-Cox transformations are given by

$$y_t = \begin{cases} \exp(w_t), & \lambda = 0; \\ (\lambda w_t + 1)^{1/\lambda}, & \lambda \neq 0. \end{cases}$$

# Back-transformation

We must reverse the transformation (or *back-transform*) to obtain forecasts on the original scale. The reverse Box-Cox transformations are given by

$$y_t = \begin{cases} \exp(w_t), & \lambda = 0; \\ (\lambda w_t + 1)^{1/\lambda}, & \lambda \neq 0. \end{cases}$$

```
lam <- BoxCox.lambda(a10) # = 0.131  
fit <- ets(a10, additive=TRUE, lambda=lam)  
plot(forecast(fit))  
plot(forecast(fit), include=60)
```



# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting**
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`.

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`.
- It does not allow re-fitting a model to new data.

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`.
- It does not allow re-fitting a model to new data.
- Use the `Arima()` function in the **forecast** package which acts as a wrapper to `arima()`.

# R functions

- The `arima()` function in the **stats** package provides seasonal and non-seasonal ARIMA model estimation including covariates.
- However, it does not allow a constant unless the model is stationary
- It does not return everything required for `forecast()`.
- It does not allow re-fitting a model to new data.
- Use the `Arima()` function in the **forecast** package which acts as a wrapper to `arima()`.
- Or use the `auto.arima()` function in the **forecast** package.

# R functions

```
auto.arima(x, d = NA, D = NA, max.p = 5, max.q = 5,  
  max.P = 2, max.Q = 2, max.order = 5,  
  start.p = 2, start.q = 2,  
  start.P = 1, start.Q = 1,  
  stationary = FALSE, ic = c("aic", "aicc", "bic"),  
  stepwise = TRUE, trace = FALSE,  
  approximation = (length(x)>100 | frequency(x)>12),  
  xreg = NULL, test = c("kpss", "adf", "pp"),  
  seasonal.test = c("ocsb", "ch"),  
  allowdrift = TRUE, lambda = NULL)
```



# R functions

```
auto.arima(x, d = NA, D = NA, max.p = 5, max.q = 5,  
  max.P = 2, max.Q = 2, max.order = 5,  
  start.p = 2, start.q = 2,  
  start.P = 1, start.Q = 1,  
  stationary = FALSE, ic = c("aic", "aicc", "bic"),  
  stepwise = TRUE, trace = FALSE,  
  approximation = (length(x)>100 | frequency(x)>12),  
  xreg = NULL, test = c("kpss", "adf", "pp"),  
  seasonal.test = c("ocsb", "ch"),  
  allowdrift = TRUE, lambda = NULL)
```

Automatic ARIMA algorithm due to Hyndman and Khandakar (JSS, 2008).

# ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.

# ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.

# ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.

# ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models have no exponential smoothing counterparts.

# ARIMA vs ETS

- Myth that ARIMA models are more general than exponential smoothing.
- Linear exponential smoothing models all special cases of ARIMA models.
- Non-linear exponential smoothing models have no equivalent ARIMA counterparts.
- Many ARIMA models have no exponential smoothing counterparts.
- ETS models all non-stationary. Models with seasonality or non-damped trend (or both) have two unit roots; all other models have one unit root.

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality**
- 8 `forecast()` function
- 9 Time series cross-validation

# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24



# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24
- `Arima()` has maximum period 350, but usually runs out of memory if period  $> 200$ .

# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24
- `Arima()` has maximum period 350, but usually runs out of memory if period  $> 200$ .
- `stl()` allows a decomposition of any frequency.

# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24
- `Arima()` has maximum period 350, but usually runs out of memory if period  $> 200$ .
- `stl()` allows a decomposition of any frequency.

# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24
- `Arima()` has maximum period 350, but usually runs out of memory if period  $> 200$ .
- `stl()` allows a decomposition of any frequency.

## Multiple seasonal periods

- `dshw()` will allow two seasonal periods.

# Difficult seasonality

## High frequency data

- `ets()` has maximum period 24
- `Arima()` has maximum period 350, but usually runs out of memory if period  $> 200$ .
- `stl()` allows a decomposition of any frequency.

## Multiple seasonal periods

- `dshw()` will allow two seasonal periods.
- Some new functions coming soon!

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 forecast() function**
- 9 Time series cross-validation

# forecast() function

- Takes a time series or time series model as its main argument

# forecast() function

- Takes a time series or time series model as its main argument
- Methods for objects of class `ts`, `ets`, `Arima`, `ar`, `HoltWinters`, `fracdiff`, `StructTS`, `stl`, etc.



# forecast() function

- Takes a time series or time series model as its main argument
- Methods for objects of class `ts`, `ets`, `Arima`, `ar`, `HoltWinters`, `fracdiff`, `StructTS`, `stl`, etc.
- Output as class `forecast`.

# forecast() function

- Takes a time series or time series model as its main argument
- Methods for objects of class `ts`, `ets`, `Arima`, `ar`, `HoltWinters`, `fracdiff`, `StructTS`, `stl`, etc.
- Output as class `forecast`.
- If first argument is class `ts`, returns forecasts from automatic ETS algorithm if non-seasonal or seasonal period is less than 13. Otherwise uses `stlf()`.

# forecast package

## > forecast(a10)

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jul 2008	24.09965	22.27328	25.91349	21.17243	26.95919
Aug 2008	24.04731	22.14468	25.99928	21.13311	27.06038
Sep 2008	24.30525	22.32954	26.33671	21.25731	27.46929
Oct 2008	25.92093	23.70871	28.15997	22.45784	29.31921
Nov 2008	26.82656	24.50077	29.29492	23.30437	30.61327
Dec 2008	31.24163	28.40841	34.08156	27.04162	35.87970
Jan 2009	33.48664	30.33970	36.69329	28.78131	38.33964
Feb 2009	20.21047	18.20746	22.25788	17.18718	23.39883
Mar 2009	22.13953	19.91242	24.39955	18.83605	25.63848
Apr 2009	22.43394	20.05864	24.83489	18.95948	26.20000
May 2009	24.03782	21.42388	26.65545	20.12685	28.21718
Jun 2009	23.79650	21.09909	26.54213	19.76471	28.14571
Jul 2009	26.03920	23.05168	29.07854	21.48762	30.95417
Aug 2009	25.94239	22.81384	29.24078	21.27523	31.01868
Sep 2009	26.18084	23.01920	29.62207	21.32134	31.56815

# Outline

- 1 Time series in R
- 2 Some simple forecasting methods
- 3 Measuring forecast accuracy
- 4 Exponential smoothing
- 5 Box-Cox transformations
- 6 ARIMA forecasting
- 7 Difficult seasonality
- 8 `forecast()` function
- 9 Time series cross-validation**

# Cross-validation

## Standard cross-validation

A more sophisticated version of training/test sets.

- Select one observation for test set, and use remaining observations in training set. Compute error on test observation.

# Cross-validation

## Standard cross-validation

A more sophisticated version of training/test sets.

- Select one observation for test set, and use remaining observations in training set. Compute error on test observation.
- Repeat using each possible observation as the test set.

# Cross-validation

## Standard cross-validation

A more sophisticated version of training/test sets.

- Select one observation for test set, and use remaining observations in training set. Compute error on test observation.
- Repeat using each possible observation as the test set.
- Compute accuracy measure over all errors.

# Cross-validation

## Standard cross-validation

A more sophisticated version of training/test sets.

- Select one observation for test set, and use remaining observations in training set. Compute error on test observation.
- Repeat using each possible observation as the test set.
- Compute accuracy measure over all errors.
- Does not work for time series because we cannot use future observations to build a model.



# Time series cross-validation

Assume  $k$  is the minimum number of observations for a training set.

- Select observation  $k + i + 1$  for test set, and use observations at times  $1, 2, \dots, k + i$  to estimate model. Compute error on forecast for time  $k + i$ .

# Time series cross-validation

Assume  $k$  is the minimum number of observations for a training set.

- Select observation  $k + i + 1$  for test set, and use observations at times  $1, 2, \dots, k + i$  to estimate model. Compute error on forecast for time  $k + i$ .
- Repeat for  $i = 0, 1, \dots, n - k - 1$  where  $n$  is total number of observations.

# Time series cross-validation

Assume  $k$  is the minimum number of observations for a training set.

- Select observation  $k + i + 1$  for test set, and use observations at times  $1, 2, \dots, k + i$  to estimate model. Compute error on forecast for time  $k + i$ .
- Repeat for  $i = 0, 1, \dots, n - k - 1$  where  $n$  is total number of observations.
- Compute accuracy measure over all errors.

# Time series cross-validation

Assume  $k$  is the minimum number of observations for a training set.

- Select observation  $k + i + 1$  for test set, and use observations at times  $1, 2, \dots, k + i$  to estimate model. Compute error on forecast for time  $k + i$ .
- Repeat for  $i = 0, 1, \dots, n - k - 1$  where  $n$  is total number of observations.
- Compute accuracy measure over all errors.

# Time series cross-validation


Assume  $k$  is the minimum number of observations for a training set.

- Select observation  $k + i + 1$  for test set, and use observations at times  $1, 2, \dots, k + i$  to estimate model. Compute error on forecast for time  $k + i$ .
- Repeat for  $i = 0, 1, \dots, n - k - 1$  where  $n$  is total number of observations.
- Compute accuracy measure over all errors.

Also called **rolling forecasting origin** because the origin ( $k + i - 1$ ) at which forecast is based rolls forward in time.


# Example: Pharmaceutical sales

Antidiabetic drug sales



# Example: Pharmaceutical sales

**Log Antidiabetic drug sales**



# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.



# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
- 2 **ARIMA model applied to log data**

# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
- 2 ARIMA model applied to log data
- 3 **ETS model applied to original data**

# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
- 2 ARIMA model applied to log data
- 3 ETS model applied to original data

# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
- 2 ARIMA model applied to log data
- 3 ETS model applied to original data

● Set  $k = 48$  as minimum training set.

# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
  - 2 ARIMA model applied to log data
  - 3 ETS model applied to original data
- Set  $k = 48$  as minimum training set.
  - Forecast 12 steps ahead based on data to time  $k + i - 1$  for  $i = 1, 2, \dots, 156$ .

# Example: Pharmaceutical sales

## Which of these models is best?

- 1 Linear model with trend and seasonal dummies applied to log data.
  - 2 ARIMA model applied to log data
  - 3 ETS model applied to original data
- Set  $k = 48$  as minimum training set.
  - Forecast 12 steps ahead based on data to time  $k + i - 1$  for  $i = 1, 2, \dots, 156$ .
  - Compare MAE values for each forecast horizon.

# Example: Pharmaceutical sales

```
k <- 48
n <- length(a10)
mae1 <- mae2 <- mae3 <- matrix(NA,n-k-1,12)
for(i in 1:(n-k-1))
{
  xshort <- window(a10,end=1995+5/12+i/12)
  xnext <- window(a10,start=1995+(6+i)/12,end=1996+(5+i)/12)
  fit1 <- tslm(xshort ~ trend + season, lambda=0)
  fcast1 <- forecast(fit1,h=12)
  fit2 <- auto.arima(xshort, lambda=0)
  fcast2 <- forecast(fit2,h=12)
  fit3 <- ets(xshort)
  fcast3 <- forecast(fit3,h=12)
  mae1[i,] <- c(abs(fcast1$mean-xnext),rep(NA,12-length(xnext)))
  mae2[i,] <- c(abs(fcast2$mean-xnext),rep(NA,12-length(xnext)))
  mae3[i,] <- c(abs(fcast3$mean-xnext),rep(NA,12-length(xnext)))
}

plot(1:12,colSums(mae3,na.rm=TRUE),type="l",col=4,xlab="horizon",ylab="MAE")
lines(1:12,colSums(mae2,na.rm=TRUE),type="l",col=3)
lines(1:12,colSums(mae1,na.rm=TRUE),type="l",col=2)
legend("topleft",legend=c("LM","ARIMA","ETS"),col=2:4,lty=1)
```

# Example: Pharmaceutical sales

